

Jacob Smith

📍 Greenville, SC ✉ jacob@smithjs.org in [linkedin.com/in/jacobthemyth](https://www.linkedin.com/in/jacobthemyth) 📄 github.com/jacobthemyth

Professional Summary

Platform Engineer / Site Reliability Engineer with extensive experience optimizing performance-critical applications and implementing infrastructure-as-code and CI/CD pipelines. Motivated by developing intuitive abstractions that allow developers to work effectively with complex systems.

Languages and technologies

- **Languages:** Ruby, Go, Bash, Lua (for Redis and nginx), SQL, C, JavaScript (Node.js), Python
 - **Distributed Systems & Databases:** PostgreSQL, Redis, Kafka, Cassandra, Elasticsearch, S3, DynamoDB
 - **DevOps:** AWS, GCP, Kubernetes, ArgoCD, GitHub Actions, Terraform, Kustomize
-

Work Experience

- | | | |
|---|---|--------------------------------------|
| DevOps Consultant | Freelance (Remote) | October 2023 - Present |
| <ul style="list-style-type: none">• Refactored an untested and brittle C/Ruby integration based on SWIG, enabling Ruby developers to make API changes without C knowledge while maintaining runtime safety and performance• Resolved a catastrophic outage by identifying and resolving virtual machine I/O bottlenecks, optimizing workload distribution across virtual block devices to account for network latency overhead• Authored detailed runbooks and automated workflows to ensure long-term system resiliency | | |
| Staff Engineer | Kajabi (Remote) | July 2021 - July 2023 |
| <ul style="list-style-type: none">• Provided technical leadership for 30+ staff across Production Engineering, UX, Quality Engineering, and Security & Risk• Guided 9 cross-functional teams in data systems design and implementation, maintainable abstractions, and performance• Mentored engineers in application performance and designing resilient interfaces to distributed systems like Kafka and DynamoDB• Designed content-addressable storage for rendering end user templates, significantly reducing memory usage across processes | | |
| Staff Production Engineer | Kajabi (Remote) | December 2020 - July 2021 |
| <ul style="list-style-type: none">• Rebuilt the CI/CD pipeline for a large Rails application using Docker and BuildKit, reducing median build time from 45 to 9 minutes through careful optimization of layer caching, dependency caching, conditional step execution, and parallel execution and as a side effect, allowed builds to be run locally with fully reproducible build artifacts.• Optimized Aurora PostgreSQL to handle 50,000 QPS during peak loads through configuration tuning and query optimization• Implemented database performance isolation by decoupling workloads with a foreign data wrapper (FDW), preserving functionality while preventing system-wide slowdowns | | |
| Senior Production Engineer (Tech Lead) | Kajabi (Remote) | November 2019 - December 2020 |
| <ul style="list-style-type: none">• Defined the hiring and onboarding processes for the Production Engineering team and grew the team from 2 to 7 engineers while managing 2 full time contractors.• Designed and executed a near-zero downtime migration of a Heroku PostgreSQL database to AWS Aurora PostgreSQL, leveraging kernel tuning and cache warming to reduce customer impact and proactive testing and rollback planning to ensure resilience.• Developed infrastructure-as-code, migration automations, and provided technical leadership to the Production Engineering team and the 6 product development teams for migrating application compute workloads from Heroku to AWS EKS.• Implemented libraries and services in Ruby, primarily around resilience in the face of partial system failure | | |
| Senior Software Consultant | Subcontractor for Test Double (Remote) | May 2019 - November 2019 |
| <ul style="list-style-type: none">• Worked on the Ruby VM (i.e. "MRI") in C to prototype the reliability and performance impact of interning all string literals.• Diagnosed and fixed memory problems in a Go service that used cgo to integrate with a Rust library that itself integrated with a C++ library, reducing memory usage of ~10k server instances by ~250MB each, significantly decreasing resource utilization across a Kubernetes cluster | | |